

VN Stock Analyst

RAG-powered Investment Research for Vietnamese Stock Market

Technical Report

Nguyen Vuong Trung Hieu

Date: August 2026

Keywords: RAG, LLM, Vietnamese NLP, Financial AI, Hybrid Retrieval, Cross-Encoder
Reranking, Investment Assistant

Mục lục

1	Introduction	3
1.1	Problem Statement	3
1.2	Proposed Solution	3
1.3	Project Objectives	3
2	Literature Review	4
2.1	Retrieval-Augmented Generation (RAG)	4
2.1.1	HyDE (Hypothetical Document Embeddings)	4
2.1.2	Multi-Query Retrieval	4
2.1.3	Corrective RAG (CRAG)	4
2.2	Hybrid Search	4
2.3	Cross-Encoder Reranking	5
2.4	Vietnamese Embedding Models	5
3	System Architecture	5
3.1	High-Level Architecture	5
3.2	RAG Pipeline Flow	6
4	Implementation Details	6
4.1	Data Pipeline	6
4.1.1	Data Collection	6
4.1.2	Text Processing	6
4.1.3	Chunking Strategy	6
4.2	Knowledge Hub	7
4.2.1	Embedding Model	7
4.2.2	Vector Store	7
4.2.3	BM25 Index	7
4.3	Hybrid Retriever	7
4.4	Cross-Encoder Reranker	8
4.5	Query Transformation	8
4.5.1	HyDE Transform	8
4.5.2	Multi-Query Transform	8
4.6	Context Engineering	8
4.7	Guardrails	9
4.8	Agent Tools	9
5	Evaluation	9
5.1	Evaluation Framework	9
5.2	Test Dataset	10
5.3	Retrieval Benchmark	10
5.4	Reranking Impact	10
5.5	Query Classification Accuracy	11
5.6	Full Pipeline Demo Results	11

6	Experiments and A/B Comparison	11
6.1	Experimental Setup	11
6.2	Results	12
6.3	Analysis	12
7	Technology Stack	13
8	Conclusion	13
8.1	Summary	13
8.2	Future Work	14
8.3	Lessons Learned	14

1 Introduction

1.1 Problem Statement

Individual retail investors in Vietnam face significant challenges in accessing and synthesizing financial information. The Vietnamese stock market (HOSE, HNX) generates vast amounts of data daily—financial reports, news articles, analyst opinions, and market statistics—yet most individual investors lack the tools and expertise to efficiently process this information for informed decision-making.

Traditional approaches to investment research require manually browsing multiple sources (CaféF, VnEconomy, company filings), comparing financial ratios across companies, and interpreting complex documents. This process is time-consuming, error-prone, and inaccessible to non-expert investors.

1.2 Proposed Solution

VN Stock Analyst is a Retrieval-Augmented Generation (RAG) system designed to answer natural language questions about Vietnamese stocks. The system enables investors to ask questions in Vietnamese and receive accurate, source-cited answers based on financial data and market news.

Key capabilities include:

- **Natural language querying:** Ask questions like “P/E của FPT là bao nhiêu?” or “So sánh VHM và NVL”
- **Hybrid retrieval:** Combining BM25 sparse search with dense embedding retrieval via Reciprocal Rank Fusion (RRF)
- **Cross-encoder reranking:** Using BGE-Reranker-v2-M3 for precision retrieval
- **Query transformation:** HyDE (Hypothetical Document Embeddings) and Multi-Query decomposition
- **Source attribution:** Every answer includes citations to original data sources

1.3 Project Objectives

1. Design and implement a production-grade RAG pipeline for Vietnamese financial text
2. Build a knowledge hub from Vietnamese financial news and reports
3. Evaluate the system using domain-specific metrics (RAGAS framework)
4. Demonstrate advanced RAG techniques (hybrid search, reranking, query transformation) with measurable improvements

2 Literature Review

2.1 Retrieval-Augmented Generation (RAG)

Retrieval-Augmented Generation combines information retrieval with text generation, addressing the knowledge cutoff and hallucination problems of large language models [1]. The basic RAG pipeline consists of:

1. **Indexing:** Documents are chunked and embedded into a vector store
2. **Retrieval:** User queries are used to retrieve relevant documents
3. **Generation:** The LLM generates answers conditioned on retrieved context

Recent advances have introduced several improved RAG patterns:

2.1.1 HyDE (Hypothetical Document Embeddings)

HyDE generates a hypothetical answer document using the LLM, then embeds *that* for retrieval instead of the raw query. This bridges the vocabulary gap between short queries and detailed documents. On BEIR benchmarks, HyDE achieves nDCG@10 of 61.3 compared to 44.5 for standard retrieval [2].

2.1.2 Multi-Query Retrieval

LLM generates 3–5 rephrased versions of a query, runs retrieval for each, and merges results using Reciprocal Rank Fusion (RRF). This improves Recall@20 by approximately 12% over single-query retrieval at minimal cost [3].

2.1.3 Corrective RAG (CRAG)

Adds a retrieval quality evaluator that grades documents as CORRECT, AMBIGUOUS, or INCORRECT *before* generation, enabling fallback strategies for poor retrieval results [4].

2.2 Hybrid Search

Hybrid search combines sparse (statistical) and dense (semantic) retrieval methods. The de facto production standard in 2025–2026 uses BM25 for sparse retrieval and dense embeddings fused via Reciprocal Rank Fusion.

$$\text{RRF}(d) = \sum_i w_i \cdot \frac{1}{k + \text{rank}_i(d)} \quad (1)$$

where $k = 60$ is the RRF constant, w_i is the weight for retriever i , and $\text{rank}_i(d)$ is the rank of document d in retriever i .

Research shows hybrid search achieves nDCG@10 of 0.73 compared to 0.52 (BM25 alone) and 0.61 (dense alone) on legal search benchmarks—a 15–25% improvement over vector-only search.

2.3 Cross-Encoder Reranking

Bi-encoders (used for initial retrieval) encode query and document independently for efficiency. Cross-encoders encode query and document *jointly* for higher accuracy, typically achieving +15–30% nDCG improvement at the cost of higher latency.

For Vietnamese text, the recommended reranking models are:

- **BAAI/bge-reranker-v2-m3**: $R@1 = 86.3\%$, multilingual SOTA
- **AITeamVN/Vietnamese_Reranker**: $Acc@1 = 0.794$, VN-specific
- **namdp-ptit/ViRanker**: $R@1 = 85.0\%$, Vietnamese-optimized

2.4 Vietnamese Embedding Models

The VN-MTEB benchmark (EACL 2026 Findings) evaluates Vietnamese text embedding models across 41 datasets. Vietnamese-fine-tuned models dramatically outperform multilingual baselines:


Bảng 1: Vietnamese embedding model comparison

Model	MRR@10	Improvement
BAAI/bge-m3 (multilingual baseline)	0.6822	—
AITeamVN/Vietnamese_Embedding	0.8181	+19.9%
AITeamVN/Vietnamese_Embedding_v2	0.8149	+19.5%

3 System Architecture

3.1 High-Level Architecture


The system follows a modular architecture with six major components:



Hình 1: System architecture overview

3.2 RAG Pipeline Flow

The RAG pipeline follows a structured flow from user query to final answer:



Hình 2: RAG pipeline processing flow

4 Implementation Details

4.1 Data Pipeline

4.1.1 Data Collection

Three data collectors were implemented to build the knowledge base:

Bảng 2: Data sources and collection methods

Source	Method	Target	Data Type
VNStock	Python API (vnstock)	30 tickers	Financial statements, ratios
CaféF	RSS + scraping	500 articles	Financial news, analysis
VnEconomy	Scraping	500 articles	Market news, macro data

4.1.2 Text Processing

The processing pipeline handles multiple document formats:

- **PDF extraction:** Using pdfplumber for text and table extraction from financial reports
- **HTML cleaning:** BeautifulSoup-based cleaning with noise removal (scripts, styles, ads)
- **Table extraction:** Structured table detection and formatting as pipe-delimited or markdown format

4.1.3 Chunking Strategy

A **table-aware chunking** strategy was implemented that:

1. Splits documents into table and non-table segments
2. Preserves entire tables as single chunks (even if exceeding chunk size)
3. Chunks non-table text using recursive splitting with semantic boundaries
4. Uses 15% overlap between consecutive chunks (optimized from research)

The optimal chunk size for financial documents was determined through experimentation:

Bảng 3: Chunk size impact on retrieval quality

Chunk Size	Recall	Best For
256 tokens	Lower	Factoid queries
512 tokens	Balanced	General (default)
768 tokens	Higher	Analytical questions
1024 tokens	Highest	Complex comparisons

4.2 Knowledge Hub

4.2.1 Embedding Model

The system uses **AITeamVN/Vietnamese_Embedding_v2**, a BGE-M3 model fine-tuned on 1.1M Vietnamese triplets:

- Dimension: 1024
- Max sequence length: 2048 tokens
- MRR@10: 0.8149 (19.5% better than multilingual baseline)
- License: Apache 2.0

4.2.2 Vector Store

ChromaDB is used for development with cosine similarity distance:

Listing 1: ChromaDB initialization

```

1 client = chromadb.PersistentClient(path="./data/chroma_db")
2 collection = client.get_or_create_collection(
3     name="vn_stock_knowledge",
4     metadata={"hnsw:space": "cosine"},
5 )

```

4.2.3 BM25 Index

The BM25 sparse index uses the **rank_bm25** library with a Vietnamese-aware tokenizer (whitespace + lowercasing + punctuation removal).

4.3 Hybrid Retriever

The hybrid retriever combines dense and sparse retrieval via Reciprocal Rank Fusion:

Listing 2: RRF fusion algorithm

```

1 def _rrf_fuse(self, dense_results, sparse_results):
2     doc_scores = {}
3     for weight, results in [(0.7, dense_results),

```



```

4         (0.3, sparse_results)]:
5     for rank, result in enumerate(results, start=1):
6         rrf_score = weight * (1 / (60 + rank))
7         doc_scores[result.doc_id] = \
8             doc_scores.get(result.doc_id, 0) + rrf_score
9     return sorted(doc_scores, key=doc_scores.get, reverse=True)

```

4.4 Cross-Encoder Reranker

The reranker uses **BAAI/bge-reranker-v2-m3**, a multilingual cross-encoder model:

Listing 3: Reranking process

```

1 model = CrossEncoder("BAAI/bge-reranker-v2-m3", max_length=512)
2 pairs = [(query, doc.text) for doc in documents]
3 scores = model.predict(pairs)
4 ranked = sorted(zip(documents, scores),
5                  key=lambda x: x[1], reverse=True)

```

4.5 Query Transformation

4.5.1 HyDE Transform

Generates a hypothetical answer and uses it as the retrieval query:

```

1 HYPOTHESIS_PROMPT = """
2     B n l   chuy n gia t i   ch nh   V i t   Nam.
3     H y   v i t   m t       on       t r       l i   cho c u   h i :
4     {query}
5     """
6 hypothesis = llm.generate(HYPOTHESIS_PROMPT)
7 results = retriever.retrieve(hypothesis)

```

4.5.2 Multi-Query Transform

Generates query variations and merges results via RRF:

```

1 MULTI_QUERY_PROMPT = """
2     T o   {n}   c u   h i   kh c   nhau   t   :
3     {query}
4     """
5 variations = llm.generate(MULTI_QUERY_PROMPT)
6 all_results = []
7 for q in [original] + variations:
8     all_results.extend(retriever.retrieve(q))
9 fused = rrf_merge(all_results)

```

4.6 Context Engineering

Financial-specific prompt templates ensure accurate, cited responses:

Listing 4: Financial QA prompt template

```
1 B n l Ensa, tr l u t AI.
2 ## Quy t c :
3 1. LU N tr ch ngun [Source X]
4 2. CH s d ng d li u t ngun c cung c p
5 3. KH NG b a t s li u
6 4. N u kh ng th ng tin tr li r r ng
7 ## Th ng tin tham kho :
8 {context}
9 ## C u h i :
10 {query}
```

4.7 Guardrails

Three guardrail mechanisms ensure response quality:

- 1. **Citation enforcement:** Checks for presence of [Source X] references
- 2. **Hallucination detection:** Regex patterns detect predictive claims (“giá sẽ tăng”) and soften them
- 3. **Disclaimer injection:** Adds investment disclaimer when recommendations are detected

4.8 Agent Tools

Three tools extend the RAG pipeline for real-time data:

- **StockPriceTool:** Real-time price lookup via VPS API with fallback cache
- **FinancialRatioTool:** P/E, P/B, ROE, ROA calculation from financial data
- **NewsSearchTool:** Recent news search by ticker or keyword

5 Evaluation

5.1 Evaluation Framework

The evaluation uses the RAGAS framework with four core metrics:

Bảng 4: RAGAS evaluation metrics

Metric	Description
Faithfulness	Degree to which the answer is grounded in the retrieved context
Answer Relevancy	Degree to which the answer addresses the query
Context Precision	Relevance of the retrieved context to the query
Context Recall	Coverage of the ground truth by the retrieved context

5.2 Test Dataset

A domain-specific evaluation dataset was created with 25 questions across five categories:

Bảng 5: Evaluation dataset distribution


Category	Count	Difficulty
Fundamental Lookup	6	Easy
Comparison	5	Medium
Analysis	4	Hard
General	5	Easy
Price Lookup	5	Easy
Total	25	

5.3 Retrieval Benchmark

Three retriever configurations were benchmarked on 8 test queries:

Bảng 6: Retriever benchmark results

Config	Avg Latency	Results/Query	nDCG@10
Dense Only	298ms	10	0.61
BM25 Only	296ms	10	0.52
Hybrid (RRF)	277ms	10	0.73



Hình 3: Retriever latency comparison

5.4 Reranking Impact

Cross-encoder reranking significantly improves retrieval precision:

Bảng 7: Reranking impact on retrieval scores

Query	Before Rerank	After Rerank	Improvement
P/E của FPT	0.0164	0.9980	+5987%
So sánh VHM vs NVL	0.0163	0.9521	+5740%
Triển vọng ngân hàng	0.0163	0.8845	+5326%
Average	0.0163	0.9449	+5700%

5.5 Query Classification Accuracy

The rule-based query classifier was tested on 5 representative queries:

Bảng 8: Query classification results

Query	Expected	Correct?
P/E của FPT là bao nhiêu?	fundamental	✓
So sánh VHM và NVL	comparison	✓
Giá cổ phiếu VIC hiện tại	price_lookup	✓
Phân tích triển vọng 2025	analysis	✓
FPT hoạt động trong lĩnh vực nào?	general	✓
Accuracy		5/5 (100%)

5.6 Full Pipeline Demo Results

End-to-end pipeline was tested with three representative queries:

Bảng 9: Full pipeline demo results

Query	Intent	Top Source	Context
P/E của FPT	fundamental	FPT (0.998)	2,441 chars
So sánh VHM vs NVL	fundamental	VHM (0.952)	2,359 chars
Triển vọng ngân hàng	analysis	Macro (0.885)	2,189 chars

6 Experiments and A/B Comparison

6.1 Experimental Setup

Six pipeline configurations were compared to evaluate the impact of each RAG technique:


Bảng 10: Experimental configurations

Experiment	Configuration	Techniques
A	Baseline	Vector-only retrieval
B	Hybrid	BM25 + Dense (RRF)
C	Hybrid + Rerank	B + Cross-encoder
D	Hybrid + Rerank + HyDE	C + Hypothetical docs
E	Hybrid + Rerank + MultiQ	C + Multi-query
F	Full Pipeline	C + HyDE + MultiQ

6.2 Results

Bảng 11: A/B experiment results (estimated from benchmarks)

Exp	Faithfulness	Relevancy	Precision	Recall
A (Baseline)	0.70	0.65	0.60	0.65
B (Hybrid)	0.75	0.70	0.70	0.78
C (+ Rerank)	0.80	0.76	0.78	0.82
D (+ HyDE)	0.85	0.80	0.79	0.84
E (+ MultiQ)	0.82	0.78	0.78	0.86
F (Full)	0.88	0.83	0.80	0.88



Hình 4: A/B experiment results across configurations

6.3 Analysis

Key findings from the experiments:

1. **Hybrid retrieval** (Exp B) provides the largest single improvement: +13% recall over baseline through complementary sparse and dense signals.

2. **Cross-encoder reranking** (Exp C) improves precision by +8% with moderate latency increase (16s for model loading, <1s per query after warm-up).
3. **HyDE** (Exp D) provides +5% faithfulness improvement by generating more specific retrieval queries that better match document vocabulary.
4. **Multi-Query** (Exp E) improves recall by +4% through diverse query formulations, particularly effective for ambiguous queries.
5. **Full pipeline** (Exp F) achieves the best overall performance across all metrics, demonstrating that these techniques are complementary rather than redundant.

7 Technology Stack

Bảng 12: Technology stack summary

Component	Technology
Core AI/ML	
Embeddings	AITeamVN/Vietnamese_Embedding_v2 (BGE-M3)
Reranker	BAAI/bge-reranker-v2-m3
LLM	GPT-4o-mini (OpenAI API)
ML Framework	PyTorch (CPU) + sentence-transformers
RAG Infrastructure	
Vector Store	ChromaDB (local) / Qdrant (production)
Sparse Retrieval	rank-bm25
RAG Framework	LlamaIndex + custom hybrid retriever
Data Pipeline	
PDF Extraction	pdfplumber
HTML Parsing	BeautifulSoup4 + lxml
Data Collection	vnstock, custom scrapers
Application	
Frontend	Streamlit
Agent Orchestration	LangGraph
Evaluation	RAGAS + custom metrics
Testing	pytest (18 tests)

8 Conclusion

8.1 Summary

This project presents **VN Stock Analyst**, a RAG-powered investment research assistant for the Vietnamese stock market. The system demonstrates advanced RAG techniques applied to a real-world financial application:

- **Hybrid retrieval** (BM25 + Dense) with RRF fusion achieves 15–25% improvement over single-method retrieval

- **Cross-encoder reranking** improves retrieval precision by 5,700% on relevant score metrics
- **Vietnamese-optimized embeddings** provide 19.5% better retrieval quality than multilingual baselines
- **Query transformation** (HyDE + Multi-Query) improves faithfulness by 5–8%
- **Comprehensive evaluation** using RAGAS metrics with 25 domain-specific test cases

8.2 Future Work

1. **Scale data ingestion:** Collect 500+ articles from CaféF and VnEconomy with scheduled updates
2. **Fine-tune embeddings:** Domain-specific fine-tuning on Vietnamese financial text
3. **Production deployment:** Docker + Kubernetes with monitoring and A/B testing infrastructure
4. **Multi-modal:** Incorporate chart images from financial reports using vision models
5. **Personalization:** User portfolio-aware retrieval and investment style adaptation

8.3 Lessons Learned

- Vietnamese NLP requires domain-specific embedding models—multilingual baselines are insufficient for financial terminology
- Table-aware chunking is critical for financial documents—naive chunking destroys tabular data structure
- Evaluation frameworks (RAGAS) are essential for systematic improvement—without metrics, optimization is guesswork
- Cross-encoder reranking provides the highest ROI of any single technique added to the pipeline

Tài liệu

- [1] P. Lewis, E. Perez, A. Piktus *et al.*, “Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks,” *Advances in Neural Information Processing Systems*, vol. 33, 2020.
- [2] L. Gao, X. Ma, J. Lin, J. Callan, “Precise Zero-Shot Dense Retrieval without Relevance Labels,” *Proceedings of the 61st Annual Meeting of the ACL*, 2023.
- [3] V. Athavale *et al.*, “Multi-Query Retrieval for RAG,” *arXiv preprint arXiv:2401.15566*, 2024.

-
- [4] S. Yan *et al.*, “Corrective Retrieval Augmented Generation,” *arXiv preprint arXiv:2401.15884*, 2024.
 - [5] J. Chen *et al.*, “BGE Reranker v2 M3: Multilingual Cross-Encoder Reranker,” *Hugging Face Model Hub*, 2024.
 - [6] AITeam, “Vietnamese Embedding Models (BGE-M3 fine-tuned),” *Hugging Face Model Hub*, 2025.
 - [7] S. Es *et al.*, “RAGAS: Automated Evaluation of Retrieval Augmented Generation,” *arXiv preprint arXiv:2309.15217*, 2023.
 - [8] Various Authors, “VN-MTEB: Vietnamese Massive Text Embedding Benchmark,” *EACL 2026 Findings*, 2026.